

Problem 1.1 (25 %)

Consider the following polynomial regression scheme of degree m , with the predictions $\hat{y}_{(i)}$ for a scalar input $x_{(i)}$ given by:

$$\hat{y}_{(i)} = w_0 + w_1 x_{(i)} + w_2 x_{(i)}^2 + \dots + w_m x_{(i)}^m. \quad (1.1.1)$$

We can fit this regressor to a given training dataset $T_{train} = \{x_{(i)}, y_{(i)}\}_{i=1}^n$, by minimizing the loss $J_{LS}(\mathbf{w})$ through the corresponding least squares problem:

$$\begin{bmatrix} \hat{y}_{(1)} \\ \hat{y}_{(2)} \\ \hat{y}_{(3)} \\ \vdots \\ \hat{y}_{(n)} \end{bmatrix} = \begin{bmatrix} 1 & x_{(1)} & x_{(1)}^2 & \dots & x_{(1)}^m \\ 1 & x_{(2)} & x_{(2)}^2 & \dots & x_{(2)}^m \\ 1 & x_{(3)} & x_{(3)}^2 & \dots & x_{(3)}^m \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{(n)} & x_{(n)}^2 & \dots & x_{(n)}^m \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \vdots \\ w_m \end{bmatrix} \quad (1.1.2)$$

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} J_{LS}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \|X\mathbf{w} - \mathbf{y}\|^2. \quad (1.1.3)$$

While least squares provides a weight configuration $\hat{\mathbf{w}}$ that minimizes the error on T_{train} , it lacks a dedicated scheme to control the *generalization* to unseen data.

Ridge regression is an extension to least squares that introduces an additional regularization parameter $\lambda \geq 0$. The adapted loss is then given by:

$$J_{\text{Ridge}}(\mathbf{w}) = \|X\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2. \quad (1.1.4)$$

In the following we will examine the role this *regularization* parameters plays in addressing *over- & underfitting*, by evaluating how least squares and Ridge regression perform on the test dataset $T_{test} = \{x_{(i)}, y_{(i)}\}_{i=n+1}^N$.

1.1.1 Download the files *regression_train.csv* and *regression_test.csv* from TUWEL and use *pandas*¹ to extract the data. Generate a scatterplot of T_{train} and T_{test} combined and highlight each dataset through a distinct color.

1.1.2 Implement the code to fit the weight vector $\hat{\mathbf{w}}$ to the training dataset T_{train} . I.e., minimize (1.1.3) for an arbitrary degree m . Use the trained regressor and calculate the MSE² of your predictions over T_{test} with $m = [0, 1, \dots, 10]$. Provide a stem plot showing the values over m . Which degree m leads to the smallest error?

Hint: Use the numpy function `pinv` for numerically stable matrix inversions.

1.1.3 Additionally compute the predictions \hat{y} of the regressor for inputs in the interval $x \in [0, 2]$ for all the configurations of m from Task 1.1.2. Depict the predictions as a line plot on top of the scatter plots from Task 1.1.1. What do you observe for high and low values of m ?

¹<https://pandas.pydata.org/>

²Mean Squared Error

1.1.4 Use matrix calculus rules to derive the analytical solution for the weight vector, which minimizes the Ridge regression loss given in (1.1.4). I. e., find $\hat{\mathbf{w}}$, such that $\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} J_{\text{Ridge}}(\mathbf{w})$. Use your result, to extend the code from Task 1.1.2 to Ridge regression. *Hint: Note, that the results should be equivalent to least squares for $\lambda = 0$.*

1.1.5 Fit a Ridge regressor with $m = 5$ to T_{train} and use the trained model to calculate the MSE on T_{test} . Plot the MSE over λ for $\lambda = [0, 0.1, 1, 5, 10, 100, 500, 1000]$. Again depict the predictions for $x \in [0, 2]$ on top of the scatter plots from Task 1.1.1. What do you observe for high and low values of λ ?

The role of the regularization parameter also becomes apparent when considering a probabilistic interpretation of least squares. Let's consider the following Bayesian approach

$$p(\mathbf{y}|\mathbf{w}) = \mathcal{N}(\mathbf{y}; \mathbf{X}\mathbf{w}, \mathbf{I}\sigma_n^2) \quad p(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \mathbf{0}, \mathbf{I}\sigma_w^2), \quad (1.1.5)$$

with a Gaussian likelihood and an i.i.d. zero-mean Gaussian prior for \mathbf{w} with variance σ_w^2 . Using Bayes's rule we can then optimize for $\hat{\mathbf{w}}$ by minimizing $J(\mathbf{w}) = -\log p(\mathbf{w}|\mathbf{y})$.

1.1.6 Find an analytical expression for $J(\mathbf{w})$ under the Bayesian interpretation given in (1.1.5) and compare the result to the loss function from (1.1.4). Then, repeat the step for a Laplacian prior with $p(\mathbf{w}) = \prod_{i=0}^m \frac{1}{2\pi\sigma_w} e^{-\frac{|w_i|}{\sigma_w}}$. Briefly discuss which solutions for $\hat{\mathbf{w}}$ are promoted by the different priors. *Hint: Note, that you can discard all terms not depending on \mathbf{w} .*

Problem 1.2 (25 %)

In the following, we will evaluate the use of linear least squares for a classification problem. As such, we are given a dataset $T_{data} = \{\mathbf{x}_{(i)}, y_{(i)}\}_{i=1}^N$ consisting of inputs $\mathbf{x} = [x_0, x_1]^T$ and their respective labels $y \in \{-1, 1\}$. For each input $\mathbf{x}_{(i)}$, the classifier generates a soft-label prediction $\tilde{y}_{(i)}$ according to:

$$\tilde{y} = w_0 x_0 + w_1 x_1 + w_2. \quad (1.2.1)$$

Following a least squares approach, the model can be trained by minimizing the corresponding loss function to obtain the weight estimate $\hat{\mathbf{w}}$:

$$J_{LS}(\mathbf{w}) = \|X\mathbf{w} - \mathbf{y}\|^2 = \|\tilde{\mathbf{y}} - \mathbf{y}\|^2 \quad \hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} J_{LS}(\mathbf{w}). \quad (1.2.2)$$

Subsequently, we use the following rule to transform the soft-label output $\tilde{y}_{(i)}$ into a hard-label $\hat{y}_{(i)}$:

$$\hat{y}_{(i)} = \begin{cases} 1, & \text{if } \tilde{y}_{(i)} \geq 0 \\ -1, & \text{otherwise.} \end{cases} \quad (1.2.3)$$

We can then examine the performance of the trained classifier by comparing the predictions $\hat{y}_{(i)}$ with the given labels $y_{(i)}$.

1.2.1 Download the files *data_blob_train.csv* and *data_blob_test.csv* from TUWEL and generate a scatterplot of each T_{train} and T_{test} . Highlight the respective class labels by color.

1.2.2 Fit the classifier to the training data, i.e., obtain the least squares solution for $\hat{\mathbf{w}}$ with X and \mathbf{y} from the training dataset. Provide the accuracy of your trained classifier for T_{train} and T_{test} . Further, generate a heatmap of the soft-label output \tilde{y} in the interval $x_0 \in [-3, 3]$ and $x_1 \in [-3, 3]$.

1.2.3 Find an expression for the decision surface, i.e., the line separating the two classes $\hat{y} = 1$ and $\hat{y} = -1$. Plot this line ontop of the scatter plots from Task 1.2.1.

1.2.4 Calculate the least squares error of your trained classifier for the samples s_1 & s_2 . I.e., obtain the loss $J_{LS}(\hat{\mathbf{w}})$ between the given label y and the soft-label \tilde{y} for each of the two samples. Which of the samples contributes the higher loss to the optimization?

$$s_1 = \{[-10, 10]^T, 1\} \quad s_2 = \{[0, 1.5]^T, -1\} \quad (1.2.4)$$

Now assume, that the hard-labels \hat{y} are used for the loss calculation and compare the results. I.e., again calculate $J_{LS}(\hat{\mathbf{w}})$ for both samples, but replace \tilde{y} with \hat{y} . Why is the optimization of the soft-labels via least squares not ideal?

1.2.5 Download the file *data_moon_train.csv* and *data_moon_test.csv* from TUWEL and again provide a scatterplot highlighting the respective class labels. Then repeat Task 1.2.2-1.2.3. Is a linear classifier an appropriate choice for this dataset?

Problem 1.3 (25 %)

The california housing dataset was derived from the 1990 US census and consists of the features shown in table 1.3.1. As such, each entry includes data for a distinct californian district with the median house value in 100 000 USD as the target variable.

Input: \mathbf{x}								Target: y
MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal

Table 1.3.1: California Housing Data

In the following, you will develop a regressor operating on this real world dataset. I.e., you will train a model that can predict the median house price for a given configuration of features. For this, we will use the Python machine learning library *scikit-learn*. The dataset is available via:

```
from sklearn.datasets import fetch_california_housing
```

Note, that you might want to set the following optional parameters:

```
as_frame=True, download_if_missing=False
```

1.3.1 Download the california housing dataset and collect the inputs and targets in a *pandas* dataframe. Examine the output of *.describe()* and *.corr()*. Which feature has the highest correlation with the target value?

1.3.2 Conduct basic visual data analysis by generating histograms for single features. Further, expose the relations between different features through scatter plots. Include a selection of plots in your report and briefly discuss your findings. *Hint: You might want to use the function relplot from the seaborn library.*

1.3.3 Use the pandas function *.sample(frac = 1, random_state = 1)* to shuffle the dataset and split it into training and test data. Use the first 15 000 elements of the shuffled dataset as T_{train} and the remaining ones as T_{test} .

1.3.4 Fit a Ridge regressor³ with $\alpha = 1$ to the training dataset T_{train} . Extract the weight vector and the intercept of your trained regressor and provide a list of the features sorted by their weights. Which features have a positive/negative impact on the house price? *Hint: Check the scikit-learn documentation for details.*

1.3.5 Compute the MAE⁴ in USD between the predictions \hat{y} of your model and the targets y for T_{train} and T_{test} separately. Visualize the error $(\hat{y}_i - y_i)$ through histograms for both subsets. Further generate a scatter plot of your predictions \hat{y} and the respective ground truth values y for the samples in T_{test} . What would the plot look like for a model without any error?

³from sklearn.linear_model import Ridge

⁴Mean Absolute Error

1.3.6 Implement 10-fold cross-validation and compute mean and standard deviation of your estimation error over all evaluation runs. Select MAE as the metric and make sure to use all available data by concatenating T_{train} and T_{test} . What do the results tell you about your regressor? *Hint: An additional resource on cross-validation is provided on TUWEL.*

1.3.7 Fit a random forest regressor⁵ from *scikit-learn* to the training dataset and repeat Task 1.3.5 and 1.3.6. Briefly discuss how the outcomes for the random forest regressor differ from Ridge regression.

⁵`sklearn.ensemble.RandomForestRegressor(n_estimators=50)`

Problem 1.4 (25 %)

In a previous problem we showed, that the naive least squares approach is not well suited for classification due to the lack of a squashing function. For a given dataset $T = \{\mathbf{x}_{(i)}, y_{(i)}\}_{i=1}^N$ with $y_{(i)} \in \{-1, 1\}$, the perceptron considers the classification objective:

$$\begin{aligned} \mathbf{x}_{(i)}^T \mathbf{w} + b &\geq 0 & \text{if } y_{(i)} = 1 \\ \mathbf{x}_{(i)}^T \mathbf{w} + b &< 0 & \text{if } y_{(i)} = -1, \end{aligned} \quad (1.4.1)$$

which is enforced through the following loss function:

$$\tilde{J}(\mathbf{w}, b) = \frac{1}{N} \sum_{i=1}^N \max(0, -y_{(i)}(\mathbf{x}_{(i)}^T \mathbf{w} + b)). \quad (1.4.2)$$

Due to the fact, that $\mathbf{w} = \mathbf{0}$ and $b = 0$ are trivial solutions to the minimization of (1.4.2), the following approximated loss function is typically considered in practice:

$$J(\mathbf{w}, b) = \frac{1}{N} \sum_{i=1}^N \log(1 + e^{-y_{(i)}(\mathbf{x}_{(i)}^T \mathbf{w} + b)}). \quad (1.4.3)$$

The hard-label outputs $\hat{y}_{(i)}$ of the perceptron are then calculated from the soft-label predictions $\tilde{y}_{(i)} = \mathbf{x}_{(i)}^T \mathbf{w} + b$ according to:

$$\hat{y}_{(i)} = \begin{cases} 1, & \text{if } \tilde{y}_{(i)} \geq 0 \\ -1, & \text{otherwise.} \end{cases} \quad (1.4.4)$$

1.4.1 In (1.4.3) the maximum expression from (1.4.2) is approximated by the function $f(t) = \log(1 + e^t)$. Plot $f(t)$ in the interval $t \in [-10, 10]$. For which combinations of $y_{(i)}$ and $\tilde{y}_{(i)}$ does $J(\mathbf{w}, b)$ saturate?

1.4.2 Find an analytical expression for the gradient $\nabla J(\mathbf{w}, b)$ of the approximated loss function from (1.4.3). For that, assume a generic weight vector with two entries $\mathbf{w} = [w_0, w_1]^T$.

1.4.3 Use the result from Task 1.4.2 and implement gradient descent to find a weight $\hat{\mathbf{w}} = [\hat{w}_0, \hat{w}_1]^T$ and bias \hat{b} estimate, that minimizes the loss $J(\mathbf{w}, b)$ over a given training dataset. In gradient descent the update equation for the estimates is given by:

$$\begin{bmatrix} \hat{\mathbf{w}}^{(k)} \\ \hat{b}^{(k)} \end{bmatrix} \leftarrow \begin{bmatrix} \hat{\mathbf{w}}^{(k-1)} \\ \hat{b}^{(k-1)} \end{bmatrix} - \alpha \cdot \nabla J(\hat{\mathbf{w}}^{(k-1)}, \hat{b}^{(k-1)}). \quad (1.4.5)$$

Implement the scheme in a way, that allows for generic values of $\hat{\mathbf{w}}^{(0)}$, $\hat{b}^{(0)}$, the step size α as well as the overall iteration duration K . Note, that $k = [1, \dots, K]$.

1.4.4 Use the code from Task 1.4.3 to fit your perceptron to the dataset *classification.csv* from TUWEL. For that use $\hat{\mathbf{w}}^{(0)} = [0, -2]$, $\hat{b}^{(0)} = 2$, $K = 600$ and $\alpha = 0.01$. Report the parameters $\hat{\mathbf{w}}^{(K)}$ and $\hat{b}^{(K)}$ and visualize the decision surface of your trained classifier. What is the accuracy of your classifier on the training dataset?

1.4.5 Save the history of the weights $\hat{w}_0^{(k)}$, $\hat{w}_1^{(k)}$ and the bias term $\hat{b}^{(k)}$ during the optimization and visualize them over the gradient descent iterations $k = [1, \dots, K]$. Also generate a plot of the optimization history of $J(\hat{\mathbf{w}}^{(k)}, \hat{b}^{(k)})$.

1.4.6 Finally, use your trained model to calculate the loss $J(\mathbf{w}, b)$ for each of the samples s_1 & s_2 . Why is the the given training objective superior to the least squares classifier from the previous problem?

$$s_1 = \{[-100, 0]^T, 1\} \quad s_2 = \{[-1, 0]^T, -1\} \quad (1.4.6)$$