

# Machine Learning Algorithms

## LVA 389.203

### Basics

Markus Rupp

3.9.2020



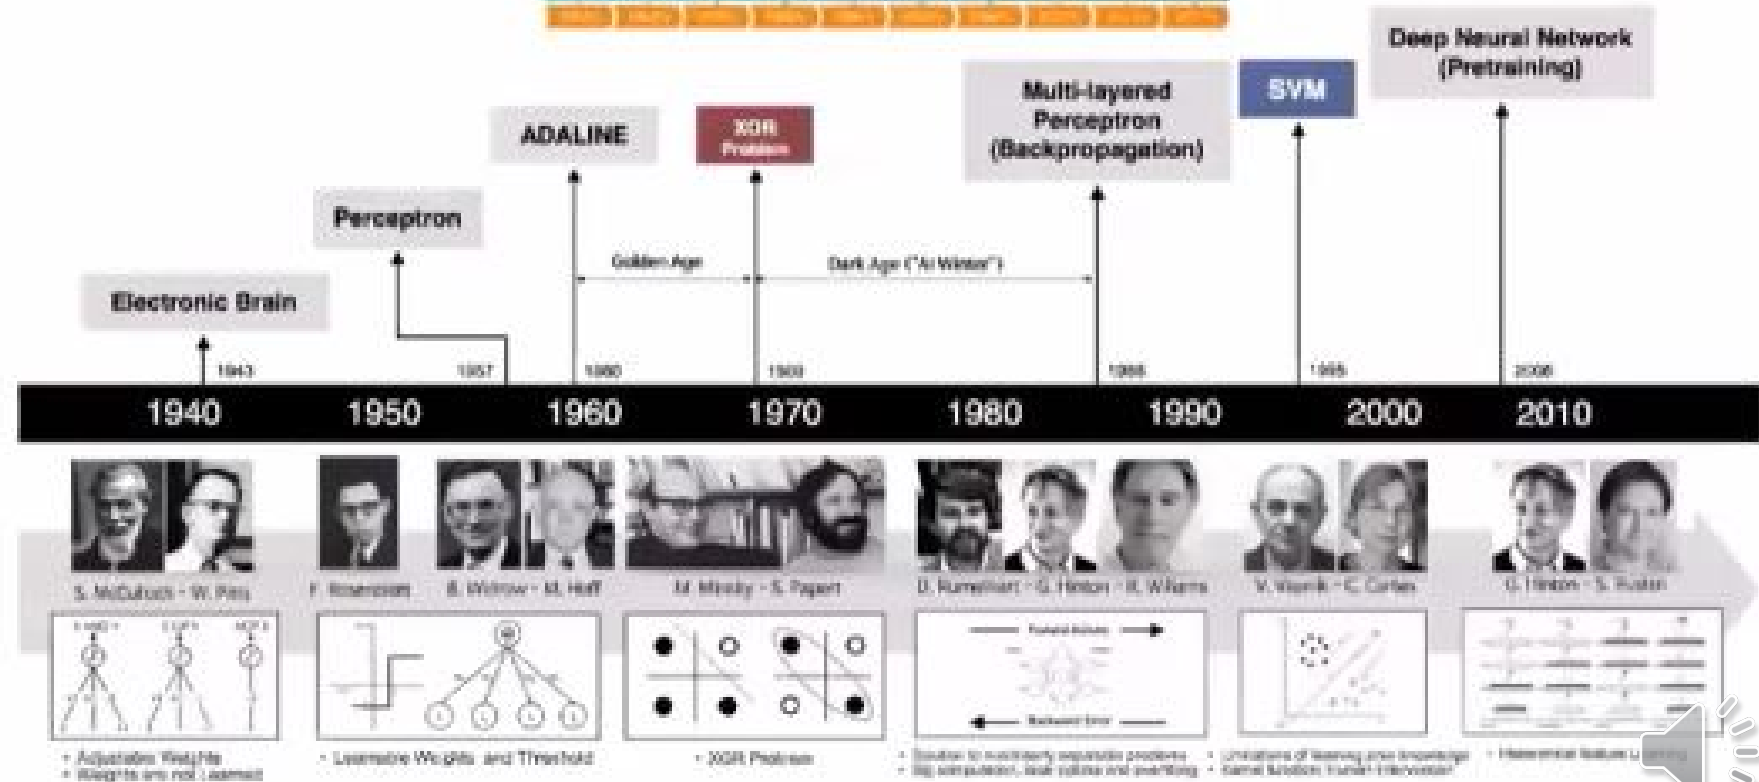
# Literature

- Watt, Borhani, Katsegallos, „Machine Learning Refined“ (1.st and 2nd edition)
- M.Rupp, „Script to Adaptive filtering course“
- Sergios Theodoridis, [Machine Learning: A Bayesian and Optimization Perspective](#), 2<sup>nd</sup> edition
- Ali Sayed, „Fundamentals of adaptive Filtering“

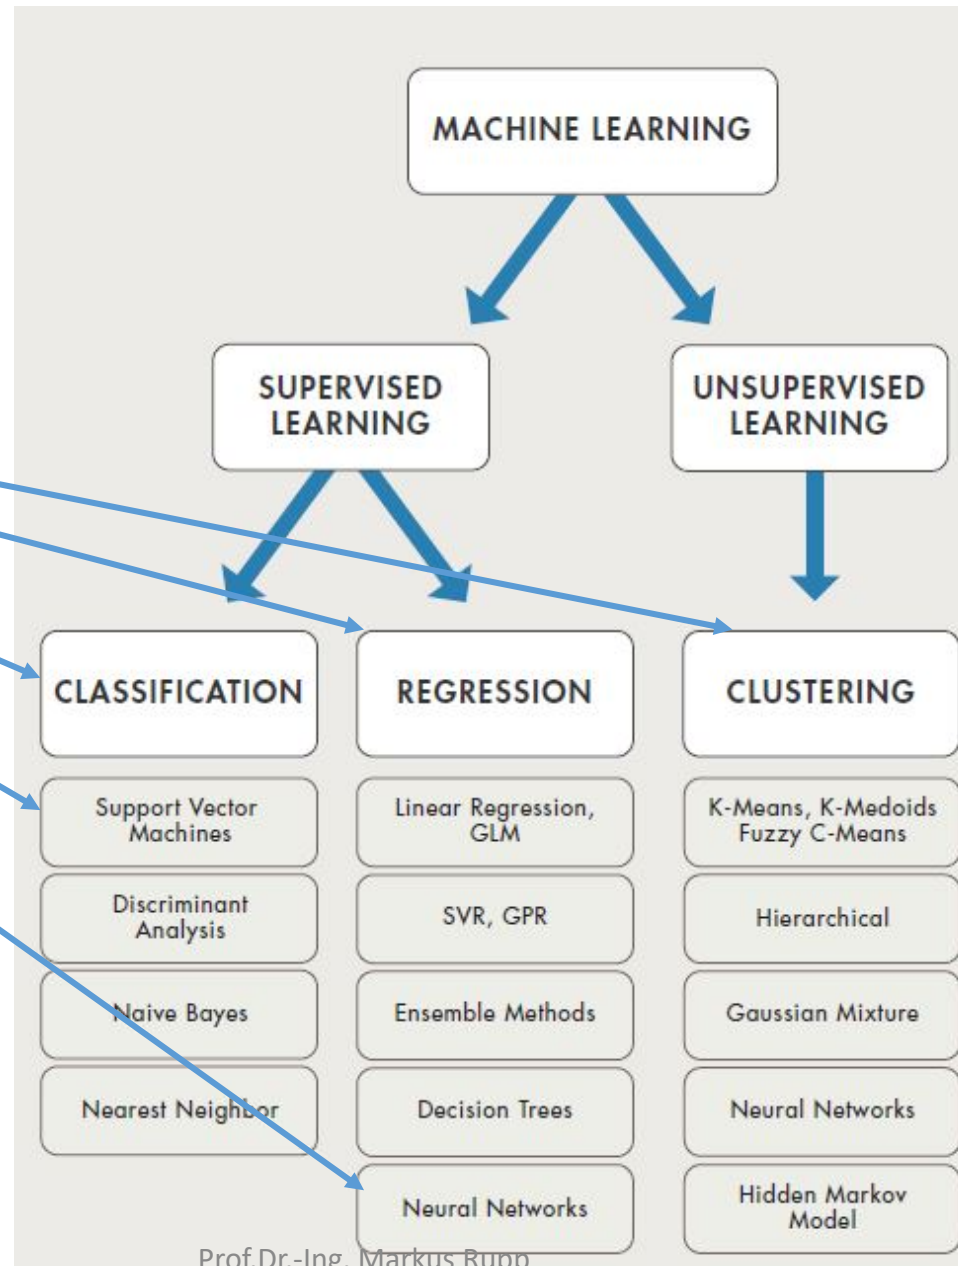


# Machine Learning (ML)

> A rapidly growing technology



This lecture



# Machine Learning

- Not entirely new topic, well researched in 80-90ies
- But some important breakthroughs in recent years.
- Interpretations based on statistics or pure data
- Strong relation to „Big Data“

Data Type	Feature Selection Task	Techniques
Sensor data	Extract signal properties from raw sensor data to create higher-level information	<b>Peak analysis</b> – perform an fft and identify dominant frequencies <b>Pulse and transition metrics</b> – derive signal characteristics such as rise time, fall time, and settling time <b>Spectral measurements</b> – plot signal power, bandwidth, mean frequency, and median frequency
Image and video data	Extract features such as edge locations, resolution, and color	<b>Bag of visual words</b> – create a histogram of local image features, such as edges, corners, and blobs <b>Histogram of oriented gradients (HOG)</b> – create a histogram of local gradient directions <b>Minimum eigenvalue algorithm</b> – detect corner locations in images <b>Edge detection</b> – identify points where the degree of brightness changes sharply
Transactional data	Calculate derived values that enhance the information in the data	<b>Timestamp decomposition</b> – break timestamps down into components such as day and month <b>Aggregate value calculation</b> – create higher-level features such as the total number of times a particular event occurred




# How do we learn?

- Learning by understanding: proof a theorem
- Learning by heart: learn a poem
- Conceptual learning: how to drive a car
- Machine learning = categorizing
  - Supervised learning: categories a priori known
  - Unsupervised learning: algorithm has to find categories




# Regression or classification?

Regression:  
Output continuous  
Can be used for  
prediction



Classification:  
Output discrete  
Can be used for  
discrimination



# (linear) Regression

- Assume a set of pairs  $(\underline{x}_i, y_i)$   $i=1,2,\dots,N$  in which we call
  - $\underline{x}_i$ : Independent variables, regressor, input
  - $y_i$ : Dependent variables, regressand, observation
- We like to find a linear combination such that



$$\begin{aligned} y_i &= c_0 + c_1 x_{i,1} + \dots + c_M x_{i,M} + v_i \\ &= \begin{bmatrix} 1, \underline{x}_i^H \end{bmatrix} \underline{c} + v_i = \underline{\hat{x}}_i^H \underline{c} + v_i \end{aligned}$$

- $v_i$  : disturbance, perturbation, error  
for all  $i=1,2,\dots,N$





# Why do we need a Bias term?



# (linear) Regression

- We like to minimize all error terms  $v_i$  in a LS manner, i.e.,

$$\begin{aligned}\min \sum_{i=1}^N |v_i|^2 &= \min \|\underline{v}\|^2 = \min_{\underline{c}} \sum_{i=1}^N \left| y_i - \hat{\underline{x}}_i^H \underline{c} \right|^2 \\ &= \min_{\underline{c}} \left\| \underline{y} - X \underline{c} \right\|^2\end{aligned}$$

$$\underline{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}; \quad X = \begin{bmatrix} \hat{\underline{x}}_1^H \\ \hat{\underline{x}}_2^H \\ \vdots \\ \hat{\underline{x}}_N^H \end{bmatrix} \in \mathbb{C}^{N \times (M+1)}$$



# (linear) Regression

- The LS solution is simply given by

$$\underline{c}_{LS} = \begin{cases} \left[ X^H X \right]^{-1} X^H \underline{y} & ; N \geq M + 1 \\ X^H \left[ X X^H \right]^{-1} \underline{y} & ; else \end{cases}$$



# (Non)-linear Regression

- By extending the basis towards nonlinear elements, it is straightforward to find a **nonlinear representation**
- For example,

$$\begin{aligned} y_i &= c_0 + c_1 x_{i,1} + \dots + c_M x_{i,M} + \\ &\quad + c_{M+1} x_{i,1}^2 + \dots + c_{2M} x_{i,M}^2 + \\ &\quad + \dots \\ &\quad + c_{KM+1} x_{i,1}^{K+1} + \dots + c_{(K+1)M} x_{i,M}^{K+1} + v_i \\ &= \begin{bmatrix} 1, \tilde{x}_i^H \end{bmatrix} \underline{c} + v_i = \hat{\underline{x}}_i^H \underline{c} + v_i \end{aligned}$$

# (Non)-linear Regression

- Note that the formulation remains linear in the coefficients and thus an ordinary LS problem.
- Strictly speaking, this is still a linear regression problem as the regressor remains linear in the coefficients.
- In practice, it is recommended to use an orthogonal polynomial family to avoid numerical problems with the Gramian.



# Classification





# Classification

In [English grammar](#), a *split infinitive* is a construction in which one or more words come between the [infinitive](#) marker *to* and the [verb](#) (as in "*to really try my best*"). Also called a *cleft infinitive*.

- - **The English-Speaking World**

"The English-speaking world may be divided into

- (1) those who neither know nor care what a [split infinitive](#) is;
- (2) those who do not know, but care very much;
- (3) those who know and condemn;
- (4) those who know and approve;
- (5) those who know and distinguish."

(H.W. Fowler and Ernest Gowers, *A Dictionary of Modern English Usage*, 2nd ed. Oxford Univ. Press, 1965)





# Classification

- Gallia est omnis divisa in partes tres, quarum unam incolunt Belgae, aliam Aquitani, tertiam qui ipsorum lingua Celtae, nostra Galli appellantur.
- J. Cesar, de Bello Gallico
- All Gaul is divided into three parts, one of which the Belgae inhabit, the Aquitani another, those who in their own language are called Celts, in our Gauls, the third.





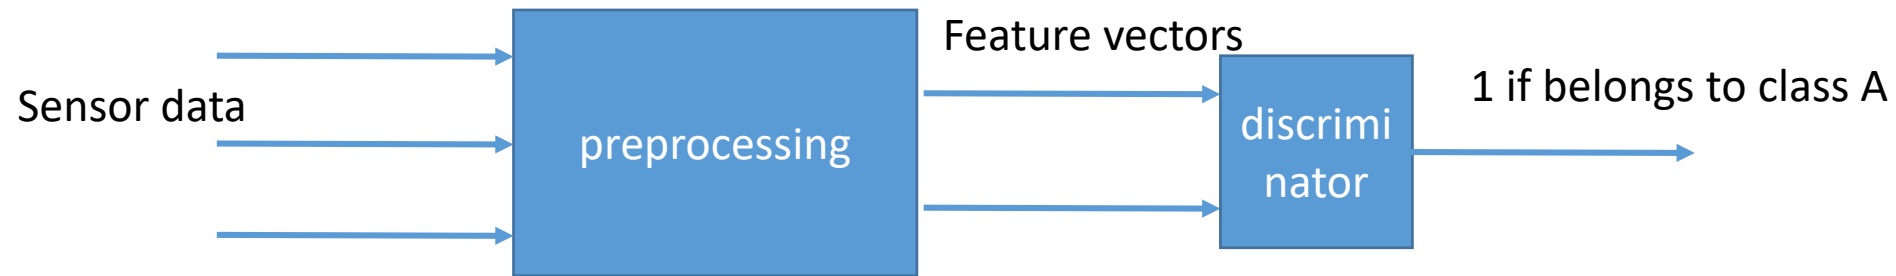
# Classification

- Assume a set of pairs  $(\underline{x}_i, y_i)$   $i=1,2,\dots,N$  in which we call
  - $\underline{x}_i$ : input vector, observation
  - $y_i$ : class, category
- Often a preprocessing of the input vectors is required to obtain a vector with desired features
  - Feature vectors:  $f(\underline{x}_i) \rightarrow \underline{f}_i$
- The problem thus reads:
  - Given a set of pairs  $(\underline{x}_i, y_i)$   $i=1,2,\dots,N$  for which a subset belongs to a class A ( $y_i=1$ ) and the remaining part not ( $y_i=-1$ ),
  - Find a hyperplane that divides the two sets




# Architectures for Classification

- Simple classification network



# Architectures for Classification

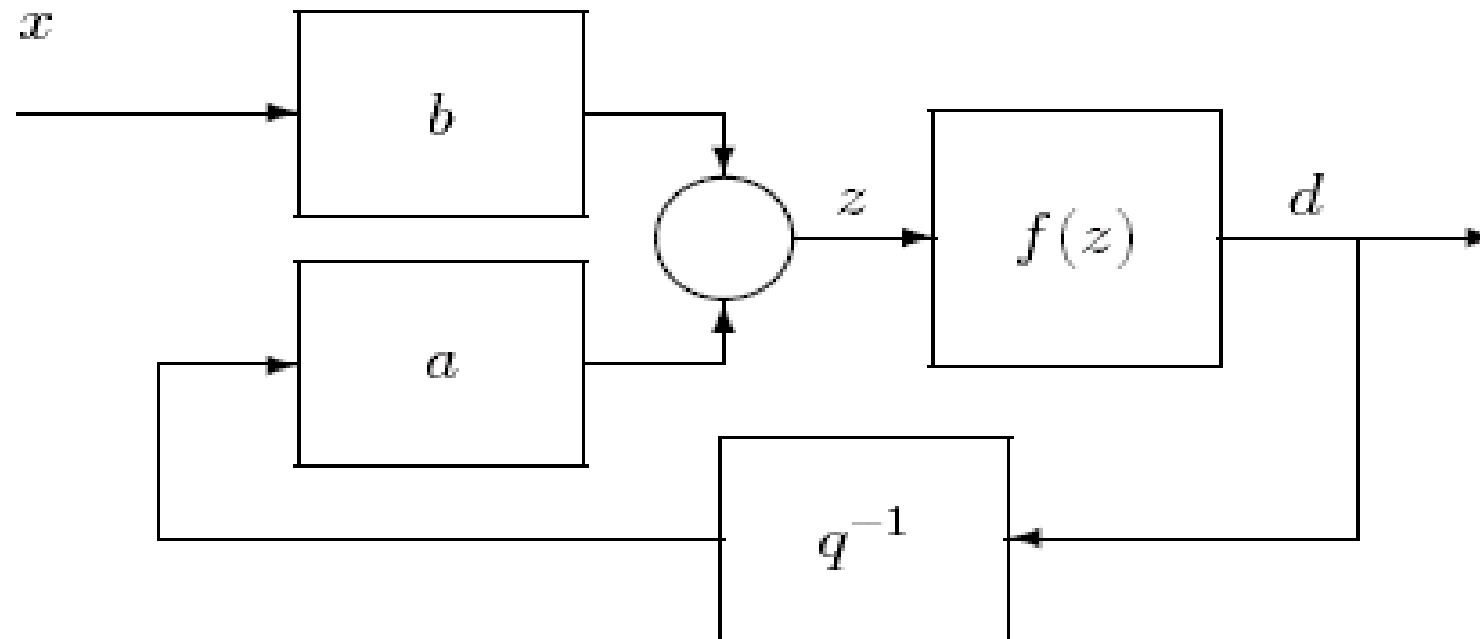
- Neural network
- Several layers of meshed and connected „neurons“



- Offers „complicated“ shapes in the feature space to distinguish

# Architectures for Classification (RNN)

Neural networks also come with feedback, in particular if time series are processed (speech detector)




# Architectures: Speed up Learning

- As learning can be very tedious and time consuming, in particular if many layers are needed, architectural structures may need to be selected depending on the problem to solve.
- While recurrent networks (RNN) are often useful in speech communications taking strong correlations into account, other architectures are useful in different contexts.
- Convolution neural networks (CNN) are particularly useful in the context of image and video understanding as predefined convolutions are able to preselect typical features as edge detections. (Le Cun 2010)

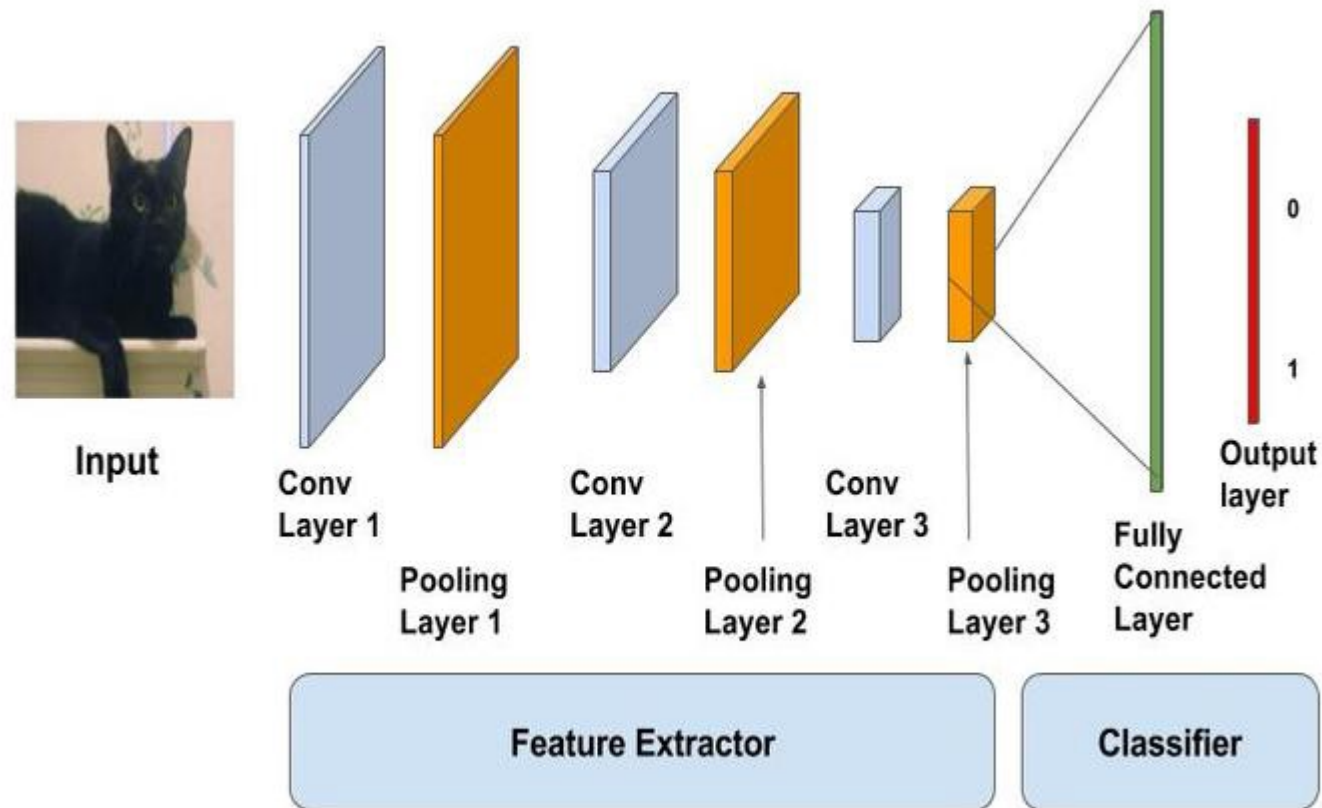


# RNN: Recurrent Neural Network

## Recurrent Neural Networks



# CNN: Convolutional Neural Network




# Architectures

- Essentially a neural network comprises of a number of layers with a selection of nonlinear functions in between. We could thus describe it parameterically.
- Specific to their associated problems it may make sense however to add preprocessors in between the layers, e.g. CNN with small convolutional units in between.
- To mimic an iterative method we can use „Unfolding“, i.e., rather than 6 iterations we add 6 layers, leading typically to deep neural networks (DNN). Recently many classic algorithms for communications have been described as DNN with specific architectural additions.



# Architecture: Autoencoding

- While most neural nets are designed to be trained on input/output pairs of data, one structure only requires an input:
- Idea is to compress data!




# Classification

- Consider a set of objects  $S$  whose members either have a certain attribute  $A$  or not. We thus want to separate the set into two subsets  $S_1$  and  $S_2$ .
- We assume each object  $s_i$  with attribute  $A$  has a certain set of features that are indicators for the desired attribute, organized in a vector  $\underline{f}_i = [f_{i1}, f_{i2}, \dots, f_{iN}]$ , while those objects that do not have such attribute have a similar vector  $\underline{g}_i$ .
- We want to find a vector  $\underline{w}$  such that the inner product  $\langle \underline{w}, \underline{f}_i \rangle$  is close to „A“ and  $\langle \underline{w}, \underline{g}_i \rangle$  to „B=not A“.



# Classification

- Consider simple example



$$y = ax + y_0$$
$$0 = [1, x, y] \begin{bmatrix} y_0 \\ a \\ -1 \end{bmatrix} = \langle \underline{\hat{f}}, \underline{w} \rangle$$

Description of line by  $\underline{w}$  is not unique

We can bring in additional constraints on  $\underline{w}$ , such as  $\min ||\underline{w}||$

$\underline{w}$  is perpendicular to the hyperplane  
 $\underline{w}$  defines the hyperplane



# Classification


- We can thus set:
$$\underline{w}_{opt} = \arg \min_{\underline{w}} \sum_{i=1}^{N_1} \left(1 - \langle \underline{w}, \underline{f}_i \rangle\right)^2 + \sum_{i=1}^{N_2} \left(-1 - \langle \underline{w}, \underline{g}_i \rangle\right)^2$$
- which results in a hyperplane that separates both sets (if they are separable).
- Note however, that such a hyperplane should also be moved to the y direction (we did the same in the regression problem), which is not possible yet.  
Therefore, we add another degree of freedom:

$$\begin{aligned} \langle [w_0, \underline{w}], [1, \underline{f}_i] \rangle &= \langle \hat{\underline{w}}, \hat{\underline{f}}_i \rangle \\ \underline{w}_{LS} &= \arg \min_{\underline{w}} \sum_{i=1}^{N_1} \left(1 - \langle \hat{\underline{w}}, \hat{\underline{f}}_i \rangle\right)^2 + \sum_{i=1}^{N_2} \left(-1 - \langle \hat{\underline{w}}, \hat{\underline{g}}_i \rangle\right)^2 \\ &= \left[ \sum_{i=1}^{N_1} \hat{\underline{f}}_i \hat{\underline{f}}_i^T + \sum_{i=1}^{N_2} \hat{\underline{g}}_i \hat{\underline{g}}_i^T \right]^{-1} \left( \sum_{i=1}^{N_1} \hat{\underline{f}}_i - \sum_{i=1}^{N_2} \hat{\underline{g}}_i \right) \end{aligned}$$




# Supervised Learning with a Linear Classifier

## Classification



# LS can fail



Thanks to  
Aggelos  
Pikrakis



# Classification

- Note that selecting „1“ and „-1“ was rather arbitrary.
- What happens if we select a and b?

$$\underline{w}_{opt} = \arg \min_{\underline{w}} \sum_{i=1}^{N_1} \left( a - \langle \underline{w}, \underline{f}_i \rangle \right)^2 + \sum_{i=1}^{N_2} \left( b - \langle \underline{w}, \underline{g}_i \rangle \right)^2$$

$$\langle [w_0, \underline{w}], [1, \underline{f}_i] \rangle = \langle \tilde{\underline{w}}, \tilde{\underline{f}}_i \rangle$$

$$\begin{aligned} \underline{w}_{LS} &= \arg \min_{\underline{w}} \sum_{i=1}^{N_1} \left( a - \langle \tilde{\underline{w}}, \tilde{\underline{f}}_i \rangle \right)^2 + \sum_{i=1}^{N_2} \left( -b - \langle \tilde{\underline{w}}, \tilde{\underline{g}}_i \rangle \right)^2 \\ &= \left[ \sum_{i=1}^{N_1} \tilde{\underline{f}}_i \tilde{\underline{f}}_i^T + \sum_{i=1}^{N_2} \tilde{\underline{g}}_i \tilde{\underline{g}}_i^T \right]^{-1} \left( b \sum_{i=1}^{N_2} \tilde{\underline{g}}_i + a \sum_{i=1}^{N_1} \tilde{\underline{f}}_i \right) \end{aligned}$$



# Classification

- Scaling  $(a,b) \rightarrow \alpha (a,b)$  results in a scaled LS solution  $\alpha \underline{w}_{LS}$ .
- As all scaled versions of  $w$  describe the same line, it does not matter.
- The distance between  $a$  and  $b$ , however, matters.



# Classification: LS or Iterative Solution?

- Our first (simple) problem allowed a formulation in LS and thus a direct solution by the computation of a pseudo inverse.
- Assuming a decent numerical posedness of the Gramian and a not too large number of variables, the ordinary LS solution would be the method of choice.
- However, often problems are formulated in nonlinear terms or at least one of the above problems occurs (in practice most times all of them occur at the same time!),
- Then **iterative schemes** are preferred.



# Iterative Schemes

- Such algorithms always work based on the following idea
  - Have a starting value (in case you have a priori knowledge, zero otherwise)
  - Improve the first estimate/approximation by a simple operation that includes an error term
  - Continue doing so until the error term becomes zero (or practically close to zero)

$$\underline{\hat{w}}_{k+1} = \underline{\hat{w}}_k + \mu \frac{\partial}{\partial \underline{w}} g(e_k^2)$$

- Advantages:
  - simple, robust, efficient in implementation
- Drawbacks:
  - may require many iterations,
  - may get stuck in a local minimum



# Iterative Schemes

- Although each iteration step often only offers a small improvement step, this property becomes very handy once new data arrive.
- Thus, with every new data pair arriving, the algorithm keeps improving its quality

$$\underline{\hat{w}}_{k+1} = \underline{\hat{w}}_k + \mu \frac{\partial}{\partial \underline{w}} g(e_k^2)$$

- However, it also begins to forget older achievements and gives the newer data more weight
- By this the iterative algorithm can perform tracking.
- If the class evolves over time, the iterative algorithm is able to follow such evolution.

# Iterative Schemes

- **Gradient based approaches** use a gradient, derived from a (quadratic) error:

$$\frac{\partial}{\partial e} g(e^2) = 2 e g'(e^2)$$

$$\frac{\partial}{\partial e} (g(e))^2 = 2 g(e) g'(e)$$

*with*  $g(0) = 0$

- **Reinforcement learning approaches** use error indicators to train the unknown weights.



# Iterative Schemes

- If the data set is fixed and we iterate along such data set, we call the algorithm **iterative or off-line algorithm**.
- If on the other hand, there is permanently new data to learn from, we call the adaptive algorithm **recursive or on-line algorithm**. Such algorithms are usually more of interest as they can potentially allow also tracking, that is to change their behaviour in case the environment changes.